

Foundations for scalable NLP-assisted formative assessment feedback

Matthew Beckman
Penn State University

April 3, 2023

Foundations for scalable NLP-assisted formative assessment feedback

Matthew Beckman
Penn State University

April 3, 2023

Two question survey before seminar (scan with mobile phone)



Figure 1: (QR Code) <https://forms.gle/hpW72fMYE1SsB19JA>

Responses to our survey?

- ① Is your lucky number odd or even?
- ② How did you describe the value of formative assessment?

Responses to our survey?

- ① Is your lucky number odd or even?
- ② How did you describe the value of formative assessment?
 - [Odd] Free text response: *write anything you like*
 - [Even] Selected response: *endorse provided options*

Motivation

- “Write-to-learn” tasks improve learning outcomes (Graham, et al., 2020)
- Critical for citizen-statisticians to communicate statistical ideas effectively (Gould, 2010)
- Continual practice with communicating improves statistical literacy and promotes retention (Basu, et al., 2013)
- A majority of U.S. undergraduates at public institutions will take at least one large enrollment STEM course (Supiano, 2022)
- Formative assessment benefits both students & instructors (Black & Wiliam, 2009; GAISE, 2016; Pearl, et al., 2012)
- *Logistics* of constructed response tasks jeopardize use in large-enrollment classes (Garfield & Ben-Zvi, 2008; Woodard & McGowan, 2012)

Goal state

Computer-assisted formative assessment feedback for short-answer tasks in large-enrollment classes, such that instructor burden is similar to small class (~30 students)

Collaborators (humans)

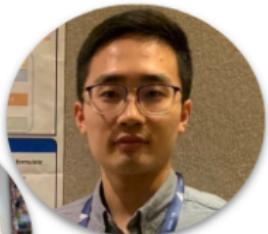
Susan Lloyd



Dennis Pearl



Zhaohui Li



Matt Beckman



Becky Passonneau



Tools (machines)

- Natural language processing (NLP) involves how computers can be programmed to analyze language elements
- NLP-assisted feedback for educational use:
 - automated short-answer grading (ASAG) from 2009
 - essays & long-answer tasks earlier
- Human-machine collaboration is a promising mechanism to assist rapid, individualized feedback at scale (Basu, 2013)
- We use Semantic Feature-Wise Transformation Relation Network (SFRN; Li, Tomar, & Passonneau, 2021)
 - back-translation data augmentation (French & Chinese)
 - can accommodate rubrics, expert solutions, or both

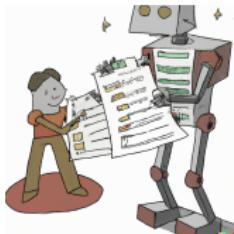


Figure 2: image created with assistance of DALL·E 2 by Open AI

SFRN Detail (Li et al., 2021)

SFRN is an end-to-end model with 3 components:

- ① encode QRA triples producing vector representations for question (Q), a possible reference (R), and student answer (A)
- ② when relation network includes multiple QRA triples, a learned feature-wise transformation network merges all relation vectors for a student answer into a single relation vector by leveraging attentions calculated by a QRA triple;
- ③ the resulting vector representation is passed as an input to a classifier (i.e., neural network)

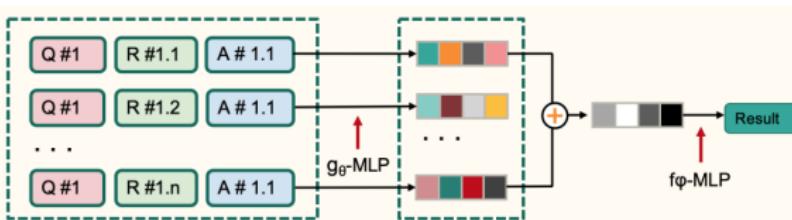
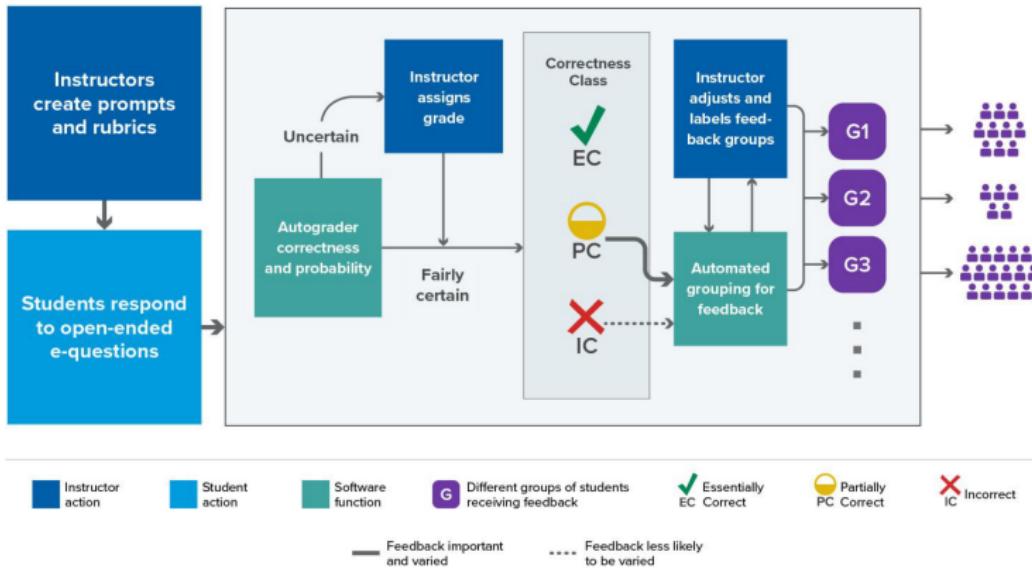


Figure 3: The g_θ MLP function computes the relation vector for each $[Q, R, A]$ triple. A set of relation vectors is combined (+) using SFT. The f_ϕ MLP function is the assessment classifier.

Project Schematic



Goal: Computer-assisted formative assessment feedback for short-answer tasks in large-enrollment classes, such that instructor burden is similar to small class (~30 students)

Research Questions

- **RQ1:** What level of agreement is achieved among trained human raters labeling (i.e., scoring) short-answer tasks?
- **RQ2:** What level of agreement is achieved between human raters and an NLP algorithm?
- **RQ3:** What sort of NLP representation leads to good clustering performance, and how does that interact with the classification algorithm?

Short Paper (ICOTS)

Lloyd, S. E., Beckman, M., Pearl, D., Passonneau, R., Li, Z., & Wang, Z. (2022). Foundations for AI-Assisted Formative Assessment Feedback for Short-Answer Tasks in Large-Enrollment Classes. In *Proceedings of the eleventh international conference on teaching statistics*. Rosario, Argentina.

Spoilers?!

- **RQ1:** What level of agreement is achieved among trained human raters labeling (i.e., scoring) short-answer tasks?
- **RQ2:** What level of agreement is achieved between human raters and an NLP algorithm?
- **RQ3:** What sort of NLP representation leads to good clustering performance, and how does that interact with the classification algorithm?

Spoilers?!

- RQ1: substantial inter-rater & intra-rater agreement
- RQ2: substantial agreement among human & NLP labeling
- RQ3: in progress, but promising

Methods (Sample)

Study utilized de-identified extant data & scoring rubrics
(Beckman, 2015)

- 6 short-answer tasks
- 1,935 students total
- 29 class sections 15 distinct institutions

Note: this sample is *not* a single large class at some institution; the available data includes introductory statistics students from many class sections at many institutions—some classes were quite small.



Figure 4: image created with assistance of DALL·E 2 by Open AI

Methods (Short-answer task)

4. Walleye is a popular type of freshwater fish native to Canada and the Northern United States. Walleye fishing takes much more than luck; better fishermen consistently catch larger fish using knowledge about proper bait, water currents, geographic features, feeding patterns of the fish, and more. Mark and his brother Dan went on a two-week fishing trip together to determine who the better Walleye fisherman is. Each brother had his own boat and similar equipment so they could each fish in different locations and move freely throughout the area. They recorded the length of each fish that was caught during the trip, in order to find out which one of them catches larger Walleye on average.

a. Should statistical inference be used to determine whether Mark or Dan is a better Walleye fisherman? Explain why statistical inference should or should not be used in this scenario.

b. Next, explain how you would determine whether Mark or Dan is a better Walleye fisherman using the data from the fishing trip. *(Be sure to give enough detail that a classmate could easily understand your approach, and how he or she would interpret the result in the context of the problem.)*

Figure 5: Sample task including a stem and two short-answer prompts.

Methods (RQ1)

- 3 human raters typical of large-enrollment instruction team
- 63 student responses in common for each *combination* of raters to quantify agreement (e.g., pairwise, consensus, etc)
- constraint: sufficient data for *intra-rater* analysis for person that had labeled 178 responses 6 years prior

Results (RQ1)

RQ1: What level of agreement is achieved among trained human raters labeling (i.e., scoring) short-answer tasks?

Comparison	Reliability
Rater A & Rater C	QWK = 0.83
Rater A & Rater D	QWK = 0.80
Rater C & Rater D	QWK = 0.79
Rater A: 2015 & 2021	QWK = 0.88
Raters A, C, & D	FK = 0.70

Reliability interpretation¹: $0.6 <$ substantial $< 0.8 <$ near perfect < 1.0

¹Viera & Garrett (2005)

Methods (RQ2)

The set of task-responses were randomly split four ways:

- 90% of data for development purposes, were partitioned according to machine-learning best practice:
 - training (72%),
 - development (9%)
 - evaluation (9%)
- 10% of data being held in reserve for more rigorous testing

SFRN was compared to other NLP algorithms for accuracy using a subset of student responses (Li et al., 2021).

- SFRN: Semantic Feature-Wise Transformation Relation Network
- LSTM: a logistic regression combined with a Long Short-Term Memory for learning vector representations

Results (RQ2)

Prerequisite-comparing machines: The SFRN algorithm achieved much higher classification accuracy than LSTM (83% vs. 72%) when judged against human consensus ratings.²

RQ2: What level of agreement is achieved between human raters and the machine (an NLP algorithm)?

Human & SFRN agreement:

Comparison	Reliability
Rater A & SFRN	QWK = 0.79
Rater C & SFRN	QWK = 0.82
Rater D & SFRN	QWK = 0.74
Raters: A, C, D, & SFRN	FK = 0.68

Reliability interpretation³: $0.6 <$ substantial $< 0.8 <$ near perfect < 1.0

²SFRN & LSTM comparison excludes instances when human labels disagree

³Viera & Garrett (2005)

Methods (RQ3)

Manual pilot of human-generated clustering

- Two reviewers independently evaluated 100 student responses that earned “partial credit” on inference tasks
- Each reviewer provided free-text feedback to each student
- Verbatim feedback captured for each reviewer and cross-tabulated for analysis.

Results (RQ3 humans)

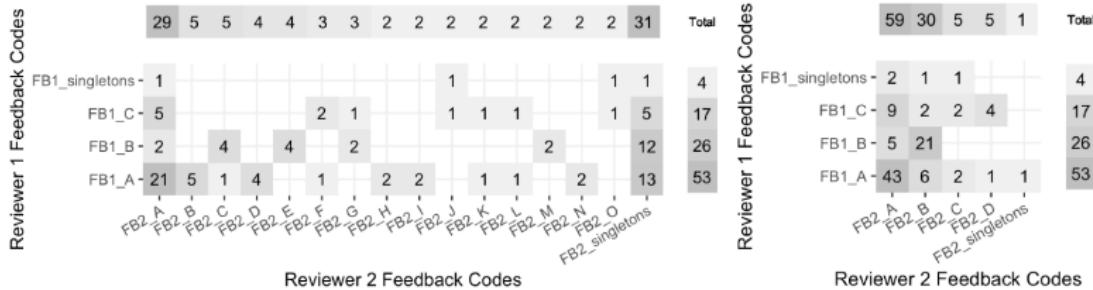


Figure 6: Cross-tabulation of feedback distribution for the two reviewers for the initial feedback (left) compared with the same analysis for the portion of feedback related to the statistical concept at issue (right).

- Reviewer 1 favored feedback on statistical concepts (only).
- Reviewer 2 provided same, plus a quote from the student
- Reviewer 2 parsed her feedback to compare her remarks related to the statistical concepts (only) with the feedback of Reviewer 1.

Results (RQ3 machines)

RQ3: What sort of NLP representation leads to good clustering performance, and how does that interact with the classification algorithm?

- SFRN learns a high-dimension ($D = 512$) vector representation on training data.
- Experiments with K-means and K-medoids clustering showed SFRN produce more consistent clusters when retrained (0.62), in comparison to LSTM *despite 8X higher dimensionality*⁴
- 3.33 correctness classes per question, on average, is reasonably consistent with human result
- Highest consistency (0.88; $D = 50$) was achieved using a matrix factorization method that produces static representations (WTMF; Guo & Diab, 2011)

⁴Consistency is measured as the ratio of all pairs of responses in a given class per question that are clustered the same way on two runs (in the same cluster, or not in the same cluster).

Discussion

- **RQ1:** Substantial agreement achieved among trained human raters provides context for further comparisons
- **RQ2:** NLP algorithm produced agreement reasonably aligned to results achieved by pairs/groups of trained human raters
- **RQ3:** Classification and clustering have competing incentives for dimensionality; Low D is better for cluster stability, High D better for classification reliability.

Limitations

- Study uses extant data from prior study collected from many classes of varying size
 - not a single large class
 - perhaps conservative due to additional variability across institutions and instructors, but will be investigated further?
- Tasks & rubrics used for pilot were developed for research use, and are likely more polished than those to be encountered “in the wild”
- Clustering performance vs semantic meaning
 - clustering is necessary, but not sufficient, for meaningful feedback
 - semantic meaning of NLP clusters not yet rigorously studied

Future Work⁵

Software development goals:

- challenge labeling algorithm (EC/PC/IC “scoring”); target linguistic diversity;
- human-in-the-loop policy to optimize conditions when algorithm should defer to human input;
- iterative instructor input to group conceptual representations
- Curse of dimensionality
 - distance between elements to be clustered increases monotonically with dimensionality (Bellman, 2003)
 - SRFN is a fairly high dimensional representation ($D = 512$)
 - For us, also perhaps a conflict of dimensionality due to tension between demands of classification step and clustering step

Field test expansion:

- field test key aspects of project CLASSIFIES in large classes
 - approx 13,000 students
 - two of the five institutions are recognized HSI's by US Dept of Education

Implications for Teaching & Research

- open questions for “what works” in formative assessment
- accumulated data made available to broader NLP community
 - this data set would be among the largest open data sources of it's kind
 - addresses barriers imposed by proprietary data sources on NLP research

References (1/3)

- ① Attali, Y., Powers, D., Freedman, M., Harrison, M., & Obetz, S. (2008). Automated Scoring of Short-Answer Open-Ended Gre® Subject Test Items. *ETS Research Report Series*, 2008(1), i–22.
- ② Basu, S., Jacobs, C., & Vanderwende, L. (2013). Powergrading: a Clustering Approach to Amplify Human Effort for Short Answer Grading. *Transactions of the Association for Computational Linguistics*, 1, 391–402. https://doi.org/10.1162/tacl_a_00236
- ③ Beckman, M. (2015). Assessment Of Cognitive Transfer Outcomes For Students Of Introductory Statistics.
<http://conservancy.umn.edu/handle/11299/175709>
- ④ Black, P., & Wiliam, D. (2009). Developing the theory of formative assessment. *Educational Assessment, Evaluation and Accountability*, 21, pp 5-31. <https://doi.org/10.1007/s11092-008-9068-5>
- ⑤ GAISE College Report ASA Revision Committee (2016). Guidelines for Assessment and Instruction in Statistics Education College Report 2016. URL: <http://www.amstat.org/education/gaise>

References (2/3)

- ⑥ Gould, R. (2010). Statistics and the Modern Student. *International Statistical Review / Revue Internationale de Statistique*, 78(2), 297–315. <https://www.jstor.org/stable/27919839>
- ⑦ Guo, W., Diab, M. (2012) Modeling Sentences in the Latent Space. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics*, pages 864–872. Association for Computational Linguistics.
- ⑧ Graham, S., Kiuhara, S. A., & MacKay, M. (2020). The Effects of Writing on Learning in Science, Social Studies, and Mathematics: A Meta-Analysis. *Review of Educational Research*, 90(2), 179–226.
- ⑨ Li, Z., Tomar, Y., & Passonneau, R. J. (2021). A Semantic Feature-Wise Transformation Relation Network for Automatic Short Answer Grading. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 6030–6040. Association for Computational Linguistics.
<https://aclanthology.org/2021.emnlp-main.487>

References (3/3)

- 10 Page, E. B. (1994). Computer Grading of Student Prose, Using Modern Concepts and Software. *The Journal of Experimental Education*, 62(2), 127–142.
- 11 Pearl, D. K., Garfield, J. B., delMas, R., Groth, R. E., Kaplan, J. J., McGowan, H., & Lee, H. S. (2012). Connecting Research to Practice in a Culture of Assessment for Introductory College-level Statistics. URL: http://www.causeweb.org/research/guidelines/ResearchReport_Doc_2012.pdf
- 12 Viera, A. J., & Garrett, J. M. (2005). Understanding interobserver agreement: the kappa statistic. *Family Medicine*, 37(5), 360–363.
- 13 Woodard, R., & McGowan, H. (2012). Redesigning a large introductory course to incorporate the GAISE guidelines. *Journal of Statistics Education*, 20(3).

Thank You

Foundations for scalable NLP-assisted formative assessment feedback

Matthew Beckman
Penn State University

April 3, 2023

Resource Page URL:

<https://mdbeckman.github.io/2023-MSU-Colloq/>

Google Photos Illustration



Same or different person?



Same



Different



Not sure

Some are harder than others. . .



Implications for Teaching & Research



Figure 7: DALLE 2 image given the prompt: “Roadmap for future educational research in the style of a retro travel poster”

Venn Diagram Fail

There was extra space on a Methods slide describing how we partitioned the sample to evaluate rater agreement... but DALLE wasn't up to the task.

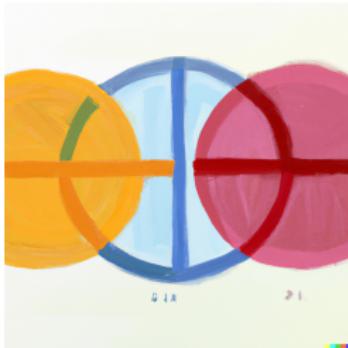


Figure 8: Me: “oil painting of Venn diagram for three intersecting sets”; DALLE 2 ... swing and a miss